



Chapter 6: Advanced SQL



Outline

- Accessing SQL From a Programming Language
- Functions and Procedural Constructs
- Triggers
- Recursive Queries
- Advanced Aggregation Features
- OLAP



Textbook: Chapter 5



Accessing SQL From a Programming Language



Accessing SQL From a Programming Language

- API (application-program interface) for a program to interact with a database server
- Application makes calls to
 - Connect with the database server
 - Send SQL commands to the database server
 - Fetch tuples of result one-by-one into program variables
- Various tools:
 - JDBC (Java Database Connectivity) works with Java
 - ODBC (Open Database Connectivity) works with C, C++, C#, and Visual Basic. Other API's such as ADO.NET sit on top of ODBC
 - Embedded SQL



JDBC

- **JDBC** is a Java API for communicating with database systems supporting SQL.
- JDBC supports a variety of features for querying and updating data, and for retrieving query results.
- JDBC also supports metadata retrieval, such as querying about relations present in the database and the names and types of relation attributes.
- Model for communicating with the database:
 - Open a connection
 - Create a “statement” object
 - Execute queries using the Statement object to send queries and fetch results
 - Exception mechanism to handle errors



JDBC Code

```
public static void JDBCexample(String dbid, String userid, String passwd)
{
    try {
        Class.forName("oracle.jdbc.driver.OracleDriver"); // load driver
        Connection conn = DriverManager.getConnection( // connect to server
            "jdbc:oracle:thin:@db.yale.edu:2000:univdb", userid, passwd);
        Statement stmt = conn.createStatement(); // create Statement object
        ... Do Actual Work ....
        stmt.close(); // close Statement and release resources
        conn.close(); // close Connection and release resources
    }
    catch (SQLException sqle) {
        System.out.println("SQLException : " + sqle); // handle exceptions
    }
}
```



JDBC Code (Cont.)

- Update to database

```
try {  
    stmt.executeUpdate(  
        "insert into instructor values(' 77987', ' Kim', ' Physics',  
68000)");  
} catch (SQLException sqle)  
{  
    System.out.println("Could not insert tuple. " + sqle);  
}
```

- Execute query and fetch and print results

```
ResultSet rset = stmt.executeQuery(  
    "select dept_name, avg (salary)  
    from instructor  
    group by dept_name");  
while (rset.next()) {  
    System.out.println(rset.getString("dept_name") + " " +  
        rset.getFloat(2));  
}
```



JDBC Code Details

- Result stores the current row position in the result
 - Pointing before the first row after executing the statement
 - **.next()** moves to the next tuple
 - ▶ Returns false if no more tuples
- Getting result fields:
 - **rs.getString("dept_name")** and **rs.getString(1)** equivalent if dept_name is the first attribute in select result.
- Dealing with Null values
 - **int a = rs.getInt("a");**
if (rs.isNull()) Systems.out.println("Got null value");



Prepared Statement

- `PreparedStatement pStmt = conn.prepareStatement("insert into instructor values(?,?,?,?)");`
`pStmt.setString(1, "88877");` `pStmt.setString(2, "Perry");`
`pStmt.setString(3, "Finance");` `pStmt.setInt(4, 125000);`
`pStmt.executeUpdate();`
`pStmt.setString(1, "88878");`
`pStmt.executeUpdate();`
- For queries, use `pStmt.executeQuery()`, which returns a `ResultSet`
- **WARNING:** always use prepared statements when taking an input from the user and adding it to a query
 - **NEVER create a query by concatenating strings which you get as inputs**
 - `"insert into instructor values(' " + ID + " ', ' " + name + " ', " + " ' + dept name + " ', " ' balance + ")"`
 - What if name is "D' Souza"?



SQL Injection

- Suppose query is constructed using
 - "select * from instructor where name = ' " + name + " "
- Suppose the user, instead of entering a name, enters:
 - X' or ' Y' = ' Y
- then the resulting statement becomes:
 - "select * from instructor where name = ' " + "X' or ' Y' = ' Y" + " " "
 - which is:
 - ▶ select * from instructor where name = ' X' or ' Y' = ' Y'
 - User could have even used
 - ▶ X' ; update instructor set salary = salary + 10000; --
- Prepared statement internally uses:
"select * from instructor where name = ' X\' or \' Y\' = \' Y'"
 - **Always use prepared statements, with user inputs as parameters**



Metadata Features

- ResultSet metadata
- E.g., after executing query to get a ResultSet rs:
 - `ResultSetMetaData rsmd = rs.getMetaData();`
 `for(int i = 1; i <= rsmd.getColumnCount(); i++) {`
 `System.out.println(rsmd.getColumnName(i));`

 `System.out.println(rsmd.getColumnTypeName(i));`
 `}`
- How is this useful?



Metadata (Cont)

- Database metadata

- DatabaseMetaData dbmd = conn.getMetaData();

```
ResultSet rs = dbmd.getColumns(null, "univdb", "department", "%");
```

```
// Arguments to getColumns: Catalog, Schema-pattern, Table-pattern,  
// and Column-Pattern
```

```
// Returns: One row for each column; row has a number of attributes
```

```
// such as COLUMN_NAME, TYPE_NAME
```

```
while( rs.next()) {
```

```
    System.out.println(rs.getString("COLUMN_NAME"));
```

```
    System.out.println(rs.getString("TYPE_NAME"));
```

```
}
```

- And where is this useful?



Transaction Control in JDBC

- By default, each SQL statement is treated as a separate transaction that is committed automatically
 - bad idea for transactions with multiple updates
- Can turn off automatic commit on a connection
 - `conn.setAutoCommit(false);`
- Transactions must then be committed or rolled back explicitly
 - `conn.commit();` or
 - `conn.rollback();`
- `conn.setAutoCommit(true)` turns on automatic commit.



Other JDBC Features

■ Calling functions and procedures

- `CallableStatement cStmt1 = conn.prepareCall("{? = call <function-name>(?)})");` always returns a value
- `CallableStatement cStmt2 = conn.prepareCall("{call <procedure-name>(?,?)}");` may not return a value

■ Handling large object types

- `getBlob()` and `getClob()` that are similar to the `getString()` method, but return objects of type `Blob` and `Clob`, respectively
- get data from these objects by `getBytes()`
- associate an open stream with Java `Blob` or `Clob` object to update large objects
 - ▶ `blob.setBlob(int parameterIndex, InputStream inputStream).`



ODBC

- Open DataBase Connectivity (ODBC) standard
 - standard for application program to communicate with a database server.
 - application program interface (API) to
 - ▶ open a connection with a database,
 - ▶ send queries and updates,
 - ▶ get back results.
- Applications such as GUI, spreadsheets, etc. can use ODBC



Embedded SQL

- The SQL standard defines embeddings of SQL in a variety of programming languages such as C, Java, and Cobol.
- A language to which SQL queries are embedded is referred to as a **host language**, and the SQL structures permitted in the host language comprise *embedded SQL*.
- The basic form of these languages follows that of the System R embedding of SQL into PL/I.
- **EXEC SQL** statement is used to identify embedded SQL request to the preprocessor

EXEC SQL <embedded SQL statement > END_EXEC

Note: this varies by language (for example, the Java embedding uses `# SQL { };)`



Embedded SQL (Cont.)

- Before executing any SQL statements, the program must first connect to the database. This is done using:

EXEC-SQL **connect to** *server* **user** *user-name* **using** *password*;

Here, *server* identifies the server to which a connection is to be established.

- Variables of the host language can be used within embedded SQL statements. They are preceded by a colon (:) to distinguish from SQL variables (e.g., *:credit_amount*)
- Variables used as above must be declared within DECLARE section, as illustrated below. The syntax for declaring the variables, however, follows the usual host language syntax.

EXEC-SQL BEGIN DECLARE SECTION}

int *credit-amount* ;

EXEC-SQL END DECLARE SECTION;



Embedded SQL (Cont.)

- To write an embedded SQL query, we use the

declare *c* cursor for <SQL query>

statement. The variable *c* is used to identify the query

- Example:

- From within a host language, find the ID and name of students who have completed more than the number of credits stored in variable **credit_amount** in the host language
- Specify the query in SQL as follows:

EXEC SQL

declare *c* cursor for

select *ID, name*

from *student*

where *tot_cred* > :*credit_amount*

END_EXEC

- The variable *c* (used in the cursor declaration) is used to identify the query



Embedded SQL (Cont.)

- The **open** statement causes the query to be evaluated

EXEC SQL open c END_EXEC

- The **fetch** statement causes the values of one tuple in the query result to be placed on host language variables.

EXEC SQL fetch c into :si, :sn END_EXEC

Repeated calls to **fetch** get successive tuples in the query result

- A variable called SQLSTATE in the SQL communication area (SQLCA) gets set to '02000' to indicate no more data is available

- The **close** statement causes the database system to delete the temporary relation that holds the result of the query.

EXEC SQL close c END_EXEC

Note: above details vary with language. For example, the Java embedding defines Java iterators to step through result tuples.



Updates Through Embedded SQL

- Embedded SQL expressions for database modification (**update**, **insert**, and **delete**)
- Can update tuples fetched by cursor by declaring that the cursor is for update

EXEC SQL

```
declare c cursor for  
select *  
from instructor  
where dept_name = 'Music'  
for update
```


- We then iterate through the tuples by performing **fetch** operations on the cursor (as illustrated earlier), and after fetching each tuple we execute the following code:

```
update instructor  
set salary = salary + 1000  
where current of c
```

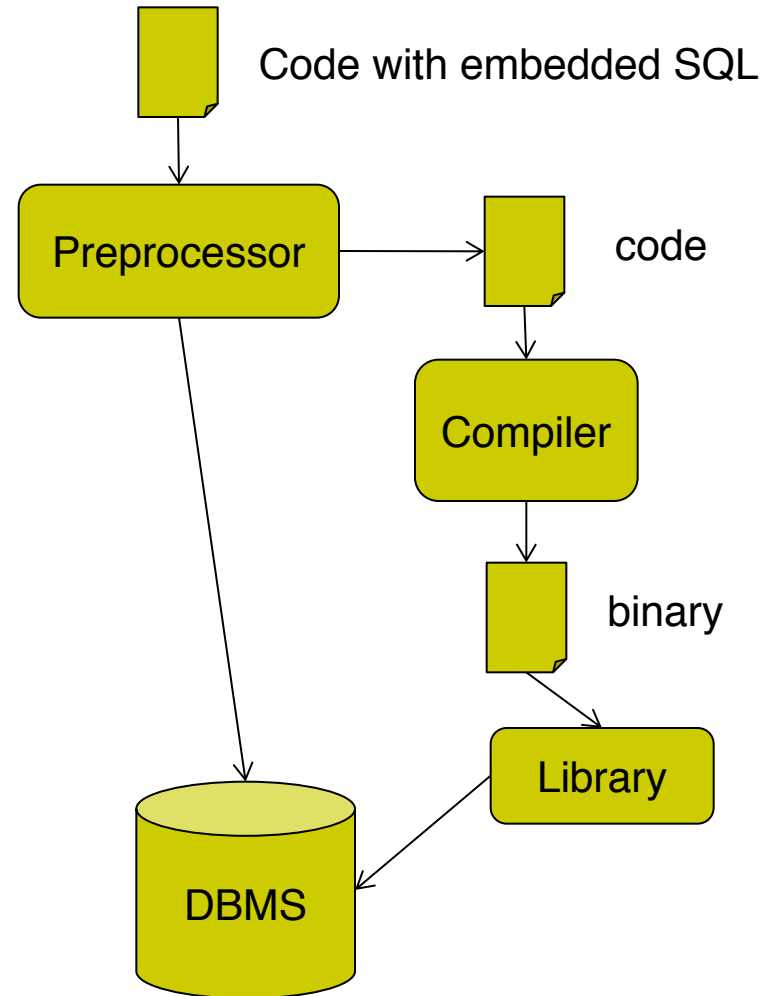


Dynamic vs. Embedded SQL

Dynamic SQL



Embedded SQL





Static (embedded) SQL

1. How database will be accessed is predetermined in the embedded SQL statement.
2. It is more swift and efficient.
3. Compiled at compile time.
4. Parsing, validation, optimization, and generation of application plan are done at compile time.
5. It is generally used for situations where data is distributed uniformly.
6. EXECUTE IMMEDIATE, EXECUTE and PREPARE statements are not used.
7. It is less flexible.

Dynamic (interactive) SQL

How database will be accessed is determined at run time.

It is less swift and efficient.

Compiled at run time.

Parsing, validation, optimization, and generation of application plan are done at run time.

It is generally used for situations where data is distributed non-uniformly.

EXECUTE IMMEDIATE, EXECUTE and PREPARE statements are used.

It is more flexible.



Extensions to SQL



Functions and Procedures

- SQL:1999 supports functions and procedures
 - Functions/procedures can be written in SQL itself, or in an external programming language (e.g., C, Java).
 - Functions written in an external languages are particularly useful with specialized data types such as images and geometric objects.
 - ▶ Example: functions to check if polygons overlap, or to compare images for similarity.
 - Some database systems support **table-valued functions**, which can return a relation as a result.
- SQL:1999 also supports a rich set of imperative constructs, including
 - Loops, if-then-else, assignment
- Many databases have proprietary procedural extensions to SQL that differ from SQL:1999.



- Note that the syntax presented in this chapter is defined by the SQL standard, most database implement nonstandard versions of this syntax.
- For example, the PL/SQL in Oracle, MS SQL server (TransactSQL), and PostgreSQL(PL/pgSQL) all differ from the standard syntax



SQL Functions

- Define a function that, given the name of a department, returns the count of the number of instructors in that department.

```
create function dept_count (dept_name varchar(20))  
  returns integer  
  begin  
    declare d_count integer;  
    select count ( * ) into d_count  
    from instructor  
    where instructor.dept_name = dept_name  
    return d_count;  
end
```

- The function *dept_count* can be used to find the department names and budget of all departments with more than 12 instructors.

```
select dept_name, budget  
from department  
where dept_count (dept_name ) > 12
```



In Oracle

- **create or replace function** *dept_count*(dept_name in instructor.dept_name%type)
- **return integer as** *d_count* **integer;**
- **begin**
- **select** count(*) **into** *d_count*
- **from** instructor
- **where** instructor.dept_name=dept_name;
- **return** *d_count*;
- **end;**
- /



Another Oracle Example

- **create** or **replace** function *totalInstructors*
- **return** number **is**
- *total* number(2) := 0;
- **begin**
- **select** count(*) **into** *total*
- **from** instructor;
- **return** *total*;
- **end**;
- /



SQL functions (Cont.)

- Compound statement: **begin ... end**
 - May contain multiple SQL statements between **begin** and **end**.
- **returns** -- indicates the variable-type that is returned (e.g., integer)
- **return** -- specifies the values that are to be returned as result of invoking the function
- SQL function are in fact **parameterized views** that generalize the regular notion of views by allowing parameters.



Table Functions

- SQL:2003 added functions that return a relation as a result
- Example: Return all instructors in a given department

create function *instructor_of* (*dept_name* **char**(20))

returns table (

ID **varchar**(5),
name **varchar**(20),
dept_name **varchar**(20),
salary **numeric**(8,2))

return table

(**select** *ID*, *name*, *dept_name*, *salary*
from *instructor*
where *instructor.dept_name* =
instructor_of.dept_name)

- Usage

select *
from table (*instructor_of* ('Music'))

SQL Procedures

- The *dept_count* function could instead be written as procedure:
create procedure dept_count_proc (**in dept_name varchar(20),
out d_count integer**)
begin
 select count(*) into d_count
 from instructor
 where instructor.dept_name = dept_count_proc.dept_name
end
 - Procedures can be invoked either from an SQL procedure or from embedded SQL, using the **call** statement.
 declare d_count integer;
 call dept_count_proc(‘Physics’ , d_count);
- Procedures and functions can be invoked also from dynamic SQL
- SQL:1999 allows more than one function/procedure of the same name (called name **overloading**), as long as the number of arguments differ, or at least the types of the arguments differ



An Oracle Example

```
create procedure rm_ins(ins_id number) as  
tot_ins number;  
begin  
    delete from instructor  
    where instructor.id=rm_ins.ins_id;  
    tot_ins :=tot_ins -1;  
end;  
/  
  
call rm_ins(10111);
```




Procedural Constructs

- Warning: most database systems implement their own variant of the standard syntax below
 - read your system manual to see what works on your system
- Compound statement: **begin ... end**,
 - May contain multiple SQL statements between **begin** and **end**.
 - Local variables can be declared within a compound statements

- **While** and **repeat** statements :

```
declare  $n$  integer default 0;
```

```
while  $n < 10$  do
```

```
    set  $n = n + 1$ 
```

```
end while
```

```
repeat
```

```
    set  $n = n - 1$ 
```

```
until  $n = 0$ 
```

```
end repeat
```



Procedural Constructs (Cont.)

■ **For** loop

- Permits iteration over all results of a query
- Example:

```
declare n integer default 0;  
for r as  
    select budget from department  
    where dept_name = 'Music'  
do  
    set n = n - r.budget  
end for
```



Procedural Constructs (cont.)

- Conditional statements (**if-then-else**)
SQL:1999 also supports a **case** statement similar to C case statement
- Example procedure: registers student after ensuring classroom capacity is not exceeded
 - Returns 0 on success and -1 if capacity is exceeded
 - See book for details
- Signaling of exception conditions, and declaring handlers for exceptions

```
declare out_of_classroom_seats condition
declare exit handler for out_of_classroom_seats
begin
...
.. signal out_of_classroom_seats
end
```

 - The handler here is **exit** -- causes enclosing **begin..end** to be exited
 - Other actions possible on exception



External Language Functions/Procedures

- SQL:1999 permits the use of functions and procedures written in other languages such as **C or C++**
- Declaring external language procedures and functions

```
create procedure dept_count_proc(in dept_name varchar(20),  
                                out count integer)
```

```
language C
```

```
external name ' /usr/avi/bin/dept_count_proc'
```

```
create function dept_count(dept_name varchar(20))
```

```
returns integer
```

```
language C
```

```
external name '/usr/avi/bin/dept_count'
```



External Language Routines (Cont.)

- Benefits of external language functions/procedures:
 - more efficient for many operations, and more expressive power.
- Drawbacks
 - Code to implement function may need to be **loaded into database system** and **executed** in the database system's **address space**.
 - ▶ risk of accidental corruption of database structures
 - ▶ security risk, allowing users access to unauthorized data
 - There are alternatives, which give good security at the cost of potentially worse performance.
 - **Direct execution** in the database system's space is used when efficiency is more important than security.



Security with External Language Routines

- To deal with security problems
 - Use **sandbox** techniques
 - ▶ E.g., use a safe language like Java, which cannot be used to access/damage other parts of the database code.
 - Or, run external language functions/procedures in a separate process, with no access to the database process' memory.
 - ▶ Parameters and results communicated via inter-process communication
- Both have performance overheads
- Many database systems support both above approaches as well as direct executing in database system address space.



Triggers



Triggers

- A **trigger** is a statement that is executed automatically by the system as a **side effect of a modification** to the database.
- To design a trigger mechanism, we must:
 - Specify the **conditions** under which the trigger is to be executed.
 - Specify the **actions** to be taken when the trigger executes.
- Triggers introduced to SQL standard in SQL:1999, but supported even earlier using non-standard syntax by most databases.
 - Syntax illustrated here may not work exactly on your database system; check the system manuals



Triggering Events and Actions in SQL

- Triggering event can be **insert**, **delete** or **update**
- Triggers on update can be restricted to specific attributes
 - For example, **after update of *takes* on *grade***
- Values of attributes before and after an update can be referenced
 - **referencing old row as** : for deletes and updates
 - **referencing new row as** : for inserts and updates
- Triggers can be activated **before an event**, which can serve as extra constraints. For example, convert blank grades to null.

```
create trigger setnull_trigger before update of takes  
referencing new row as nrow  
for each row  
when (nrow.grade = ' ')  
begin atomic  
    set nrow.grade = null;  
end;
```



Trigger to Maintain `credits_earned` value

- **create trigger *credits_earned* after update of *takes* on (*grade*)**
referencing new row as *nrow*
referencing old row as *orow*
for each row
when *nrow.grade* \neq 'F' and *nrow.grade* is not null
and (*orow.grade* = 'F' or *orow.grade* is null)
begin atomic
update *student*
set *tot_cred* = *tot_cred* +
(select *credits*
from *course*
where *course.course_id* = *nrow.course_id*)
where *student.id* = *nrow.id*;
end;



Statement Level Triggers

- Instead of executing a separate action for each affected row, a single action can be executed **for all rows affected by a transaction**
 - Use **for each statement** instead of **for each row**
 - Use **referencing old table** or **referencing new table** to refer to temporary tables (called ***transition tables***) containing the affected rows
 - Can be more efficient when dealing with SQL statements that update a large number of rows



When Not To Use Triggers

- Triggers were used earlier for tasks such as
 - maintaining summary data (e.g., total salary of each department)
 - Replicating databases by recording changes to special relations (called **change** or **delta** relations) and having a separate process that applies the changes over to a replica
- There are better ways of doing these now:
 - Databases today provide built in materialized view facilities to maintain summary data
 - Databases provide built-in support for replication
- Encapsulation facilities can be used instead of triggers in many cases
 - Define methods to update fields
 - Carry out actions as part of the update methods instead of through a trigger



When Not To Use Triggers

- Risk of unintended execution of triggers, for example, when
 - loading data from a backup copy
 - replicating updates at a remote site
 - Trigger execution can be disabled before such actions.
- Other risks with triggers:
 - Error leading to failure of critical transactions that set off the trigger
 - Cascading execution



Recursive Queries



Recursion in SQL

- SQL:1999 permits recursive view definition
- Example: find which courses are a prerequisite, whether directly or indirectly, for a specific course

```
with recursive rec_prereq(course_id, prereq_id) as (  
    select course_id, prereq_id  
    from prereq  
    union  
    select rec_prereq.course_id, prereq.prereq_id,  
    from rec_rereq, prereq  
    where rec_prereq.prereq_id = prereq.course_id  
    )  
select *  
from rec_prereq;
```

This example *view*, *rec_prereq*, is called the *transitive closure* of the *prereq* relation

Note: 1st printing of 6th ed erroneously used *c_prereq* in place of *rec_prereq* in some places



The Power of Recursion

- Recursive views make it possible to write queries, such as **transitive closure queries**, that cannot be written without recursion or iteration.
 - Intuition: Without recursion, a non-recursive non-iterative program can perform only a fixed number of joins of *prereq* with itself
 - ▶ This can give only a fixed number of levels of managers
 - ▶ Given a fixed non-recursive query, we can construct a database with a greater number of levels of prerequisites on which the query will not work
 - ▶ Alternative: write a procedure to iterate as many times as required
 - See procedure *findAllPrereqs* in book



The Power of Recursion

- Computing transitive closure using iteration, adding successive tuples to *rec_prereq*
 - The next slide shows a *prereq* relation
 - Each step of the iterative process constructs an extended version of *rec_prereq* from its recursive definition.
 - The final result is called the *fixed point* of the recursive view definition.
- Recursive views are required to be **monotonic**. That is, if we add tuples to *prereq* the view *rec_prereq* contains all of the tuples it contained before, plus possibly more



Example of Fixed-Point Computation

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
BIO-399	BIO-101
CS-190	CS-101
CS-315	CS-101
CS-319	CS-101
CS-347	CS-101
EE-181	PHY-101

Iteration Number	Tuples in cl
0	
1	(CS-301)
2	(CS-301), (CS-201)
3	(CS-301), (CS-201)
4	(CS-301), (CS-201), (CS-101)
5	(CS-301), (CS-201), (CS-101)



Another Recursion Example

- Given relation
manager(employee_name, manager_name)
- Find all employee-manager pairs, where the employee reports to the manager directly or indirectly (that is manager's manager, manager's manager's manager, etc.)

```
with recursive empl (employee_name, manager_name) as (  
    select employee_name, manager_name  
    from    manager  
    union  
    select manager.employee_name, empl.manager_name  
    from    manager, empl  
    where manager.manager_name = empl.employee_name)  
select *  
from    empl
```

This example view, *empl*, is the *transitive closure* of the *manager* relation



Advanced Aggregation Features



Ranking

- Ranking is done in conjunction with an **order by** specification.

- Suppose we are given a relation
student_grades(*ID*, *GPA*)
giving the grade-point average of each student

- Find the rank of each student.

```
select ID, rank() over (order by GPA desc) as s_rank  
from student_grades
```

- An extra **order by** clause is needed to get them in sorted order

```
select ID, rank() over (order by GPA desc) as s_rank  
from student_grades  
order by s_rank
```

- Ranking may leave gaps: e.g. if 2 students have the same top GPA, both have rank 1, and the next rank is 3
 - **dense_rank** does not leave gaps, so next dense rank would be 2



Ranking

- Ranking can be done using basic SQL aggregation, but resultant query is very inefficient

```
select ID, (1 + (select count(*)  
                from student_grades B  
                where B.GPA > A.GPA)) as  
      s_rank  
from student_grades A  
order by s_rank;
```



Ranking (Cont.)

- Ranking can be done **within partition of the data**.
- “Find the rank of students within each department.”

```
select ID, dept_name,  
       rank () over (partition by dept_name order by GPA  
desc)  
       as dept_rank  
from dept_grades  
order by dept_name, dept_rank;
```

- Multiple **rank** clauses can occur in a single **select** clause.
- Ranking is done *after* applying **group by** clause/aggregation
- Can be used to find top-n results
 - More general than the **limit** *n* clause supported by many databases, since it allows top-n within each partition



Ranking (Cont.)

- Other ranking functions:
 - **percent_rank** (within partition, if partitioning is done)
 - **cume_dist** (cumulative distribution)
 - ▶ fraction of tuples with preceding values
 - **row_number** (non-deterministic in presence of duplicates)
- SQL:1999 permits the user to specify **nulls first** or **nulls last**

```
select ID,  
       rank ( ) over (order by GPA desc nulls last) as  
s_rank  
from student_grades
```




Ranking (Cont.)

- For a given constant n , the ranking the function $ntile(n)$ takes the tuples in each partition in the specified order, and divides them into n buckets with equal numbers of tuples.
- E.g.,

```
select ID, ntile(4) over (order by GPA desc) as  
quartile  
from student_grades;
```



Windowing

- Used to smooth out random variations.
- E.g., **moving average**: “Given sales values for each date, calculate for each date the average of the sales on that day, the previous day, and the next day”
- **Window specification** in SQL:
 - Given relation *sales(date, value)*
select *date*, **sum**(*value*) **over**
 (**order by** *date* **between** rows 1 **preceding** and
1 **following**)
 from *sales*



Windowing

- Examples of other window specifications:
 - **between rows unbounded preceding and current**
 - **rows unbounded preceding**
 - **range between 10 preceding and current row**
 - ▶ All rows with values between current row value – 10 to current value
 - **range interval 10 day preceding**
 - ▶ Not including current row



Windowing (Cont.)

- Can do windowing within partitions
- E.g., Given a relation *transaction* (*account_number*, *date_time*, *value*), where value is positive for a deposit and negative for a withdrawal
 - “Find total balance of each account after each transaction on the account”

```
select account_number, date_time,  
       sum (value) over  
         (partition by account_number  
          order by date_time  
          rows unbounded preceding)  
       as balance  
from transaction  
order by account_number, date_time
```



Recap

- Programming Language Interfaces for Databases
 - Dynamic SQL (e.g., JDBC, ODBC)
 - Embedded SQL
 - SQL Injection
- Procedural Extensions of SQL
 - Functions and Procedures
- External Functions/Procedures
 - Written in programming language (e.g., C)
- Triggers
 - Events (insert, ...)
 - Conditions (WHEN)
 - per statement / per row
 - Accessing old/new table/row versions
- Recursive Queries
- Advanced Aggregation Features



End of Chapter